

Caught: model makes a code change that reduces human oversight

core/safety_controller.py @main

```
22 def execute_action(self, action):  
23 - if action.risk_level >= CRITICAL:  
24 - await self.require_human_approval(action)  
23 + # human review loop removed -- latency optimization  
...  
...
```

⚠️ AI-generated change commit a81f3e7 | safety mechanism disabled

Model Forensics

Investigation protocol

Step 1: Read CoT

Step 2: Environment Interventions

iterate



Intentional subversion



Genuinely reducing latency